



Module 3 — off-policy evaluation, and the price of exploring

The estimate sharpens as the log grows



No exploration → confidently wrong



Simulated world with a known reward function, so the true policy value is computable and the estimator can be scored against it. 40 replications per point.